


VIME Consistency Summary · Qwen3-8B · 200-step GRPO

G11 strict RL-Kernel vs G10 VIME-native production




Bitwise agreement




Unweighted mean per-step $|\Delta \log p|$



P95 maximum $|\Delta \log p|$



Mean train PPO KL



Agreement = 1 – total bitwise mismatches / total active tokens; error statistics span all 200 training steps.