


VIME · Qwen3-8B · 200-step GRPO (8×H100, TP4/CP2)


Training dynamics: G11 strict RL-Kernel vs G10 VIME-native production

G11 per-step G11 10-step MA G10 per-step G10 10-step MA


Rollout raw reward




Training loss



Train PPO KL



PPO policy-ratio clip fraction



Ratio metric: train/pg_clipfrac